Multiprocessing and Parallelism

Ways to achieve parallelism

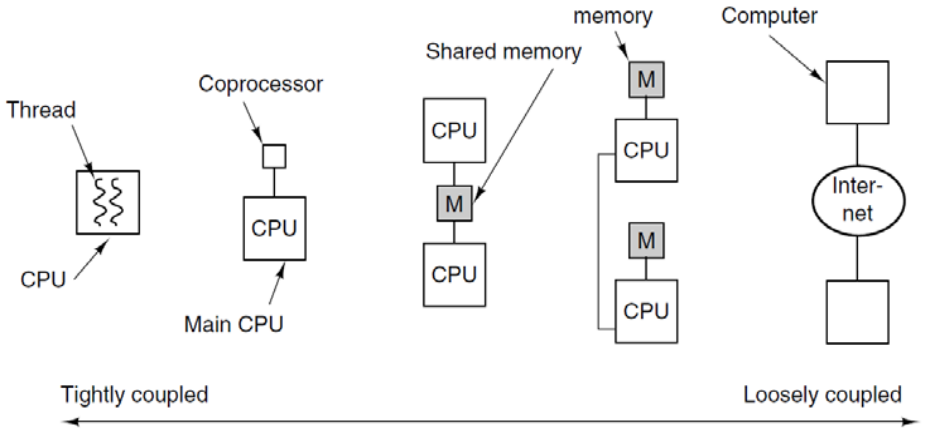

On-chip parallelism

1. **Instruction level parallelism**
   - Issue multiple instructions per clock cycle

Example: Dual-pipeline architecture

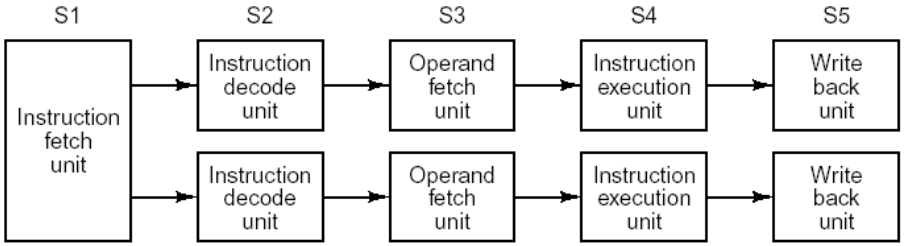

- Instruction fetch unit issues two instructions, puts each into its own pipeline
- In case of a conflict, one pipeline is executed, the other is held, and paired with next instruction

Example: Superscalar architecture

- A single pipeline with multiple functional units
- VLIW (Very Long Instruction Word)(many opcodes and sets of operands)
- Can perform 5 operations at the same time (5 opcodes + 5 pairs of operands)

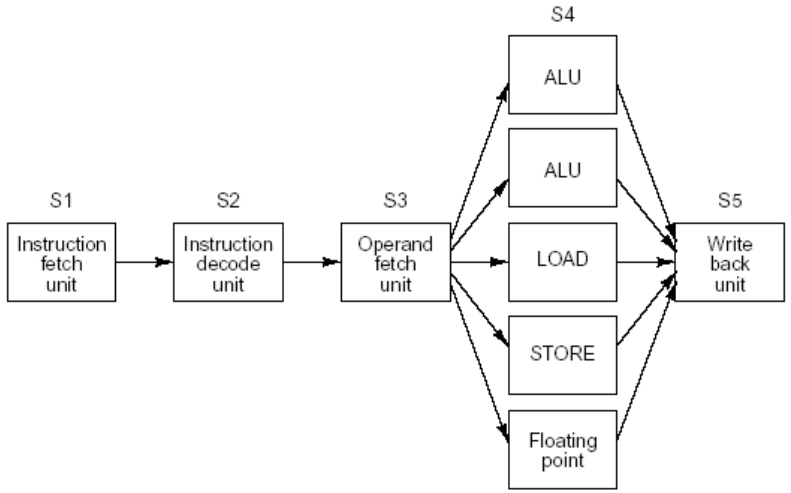

On-chip (hardware) multithreading

- In a pipelined CPU, pipeline stalls in a cache miss
- This allows the CPU to manage multiple threads at the same time in order to mask these stalls
  - Thread is a lightweight process that includes program counter, register state, stack (details will be discussed in OS course)
- Increases the utilization of a processor by switching to another thread when one is stalled

Variants of on-chip multithreading

- ✓ Fine-grained multithreading
- ✓ Coarse-grained multithreading
- ✓ Simultaneous multithreading

Fine-grained multithreading

- Switches between threads on each instruction
- Runs threads in round robin fashion with a different thread in consecutive cycles
- If an instruction stalls, subsequent instructions cannot be issued

Example:

- Assume a computer that can issue one instruction per clock cycle
- Three threads (A, B, C), 20 clock cycles
- Blank cells: dead cycles (2 cycles for miss in level 1 cache, 5 cycles for level 2, 50 cycles for last level cache)

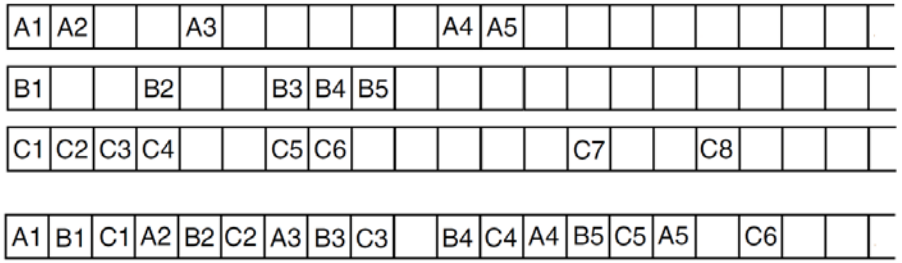

- Advantage:
  - Hide throughput losses arising from stalls
- Disadvantage:
  - Thread that is ready execute without stalls will be delayed by instructions from other threads

Coarse-grained multithreading

- A thread starts and continues to issue instructions until it encounters an expensive stall (eg: last level cache miss)
- At this point a switch occurs and the next thread starts
- Threads are run in turn in this manner

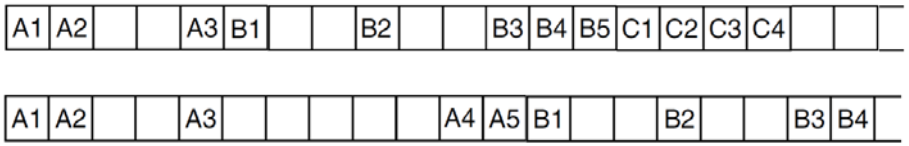

- Less efficient than fine-grained
- Unable to overcome throughput losses from shorter stalls

Fine-grained multithreading in a dual-issue CPU

- CPU can issue 2 instructions per thread in a clock cycle
- When an instruction stalls, subsequent instructions cannot be issued

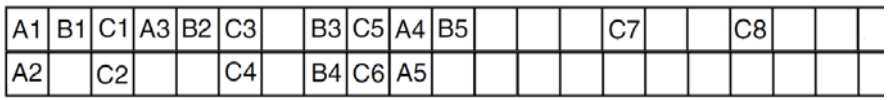

Coarse-grained multithreading in a dual-issue CPU

- CPU issues 2 instructions per thread in a clock cycle
- Threads are run in turn until one instruction stalls (larger) and then switch to the next thread in the next cycle immediately

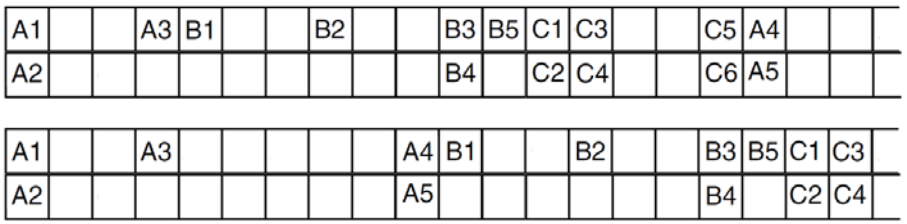

Simultaneous multithreading in a dual-issue CPU

- Refinement to coarse-grained multithreading
- A single thread is issues two instructions per clock cycle as long as it can
- When stalls, instructions are immediately taken from the next thread in sequence

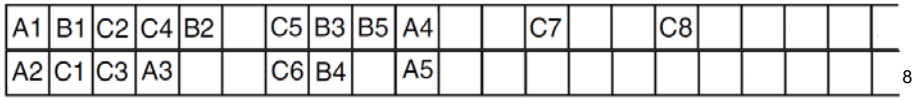

Hyper-Threading (HT)

- Allows a single processor to run two treads at once
- HT converts a single physical processor into two virtual processors
- An HT-enabled processor has 2 sets of general-purpose registers, control registers and other architecture components, but same cache, execution units and buses
- Provides 25% increase in performance for additional 5% increase in chip size for HT hardware

Single-chip multiprocessors

- Homogeneous multiprocessors on a chip
  - Having 2 or more CPUs on a single chip
  - Share same memory, last level of cache etc.
- Heterogeneous multiprocessors on a chip
  - Appears in audio-visual consumer electronics (TV, DVD players etc.)

Coprocessors

- A second specialized processor
- Coprocessors perform specialized tasks helping the main processor
- Come in as a part of the CPU package or in a plug-in board

Math coprocessor (Floating-Point Unit) (FPU)

- Provides hardware for floating-point math
- Speeds computer's operations when running software designed to use the coprocessor
- Can perform high-level mathematical operations (trigonometric functions, roots, logarithms) many times faster
- Instruction set of math chip is different from that of main CPU

GPU (Graphics Processing Unit)

- Designed to handle high-resolution graphics processing
- Contains a large number of cores with a smaller instruction set
- Rely on hardware multithreading
- Comes in as
  - •Plug-in cards (high performance)
  - •On motherboard (high end video used in laptops)
  - •In motherboard chipset (economical, shares system memory and other components, less powerful)
  - •Integrated to processor (shares system memory, less powerful)

Shared Memory Multiprocessor (SMP)

- All CPUs share a common memory
- Processors communicate through shared variables in memory
- Allows old programs to run well on parallel hardware

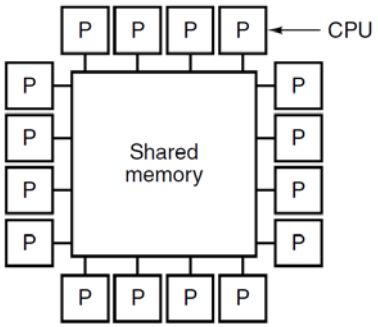

Message-Passing Multiprocessors

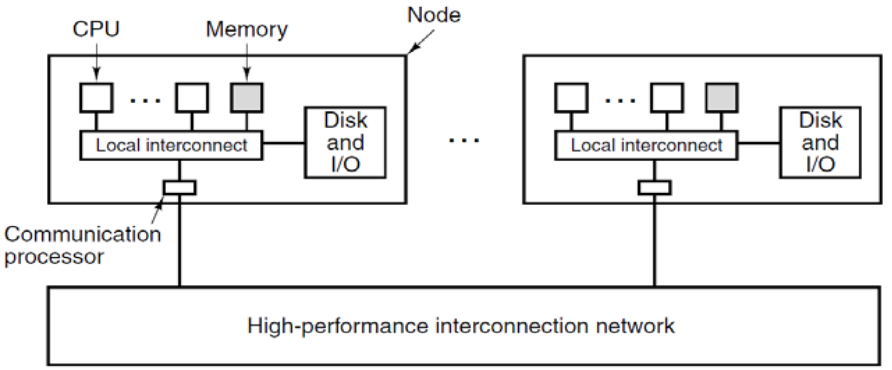

- Each node consists of
  - One or more CPUs
  - RAM shared by CPUs within the node
  - Disk and I/O devices
  - Communication processor
- Communication processors are connected by a high-speed interconnection network
- Processors communicate by passing messages