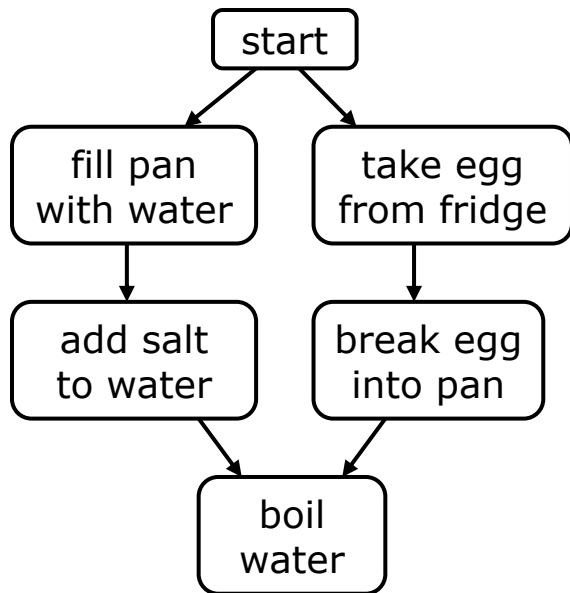


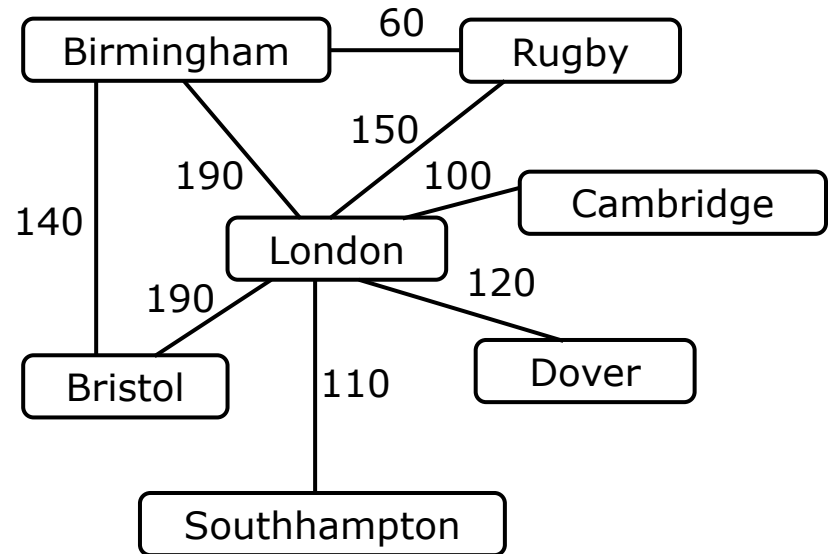
Graphs

Graph definitions

- There are two kinds of graphs: **directed graphs** (sometimes called digraphs) and **undirected graphs**



A directed graph



An undirected graph

Graph terminology I

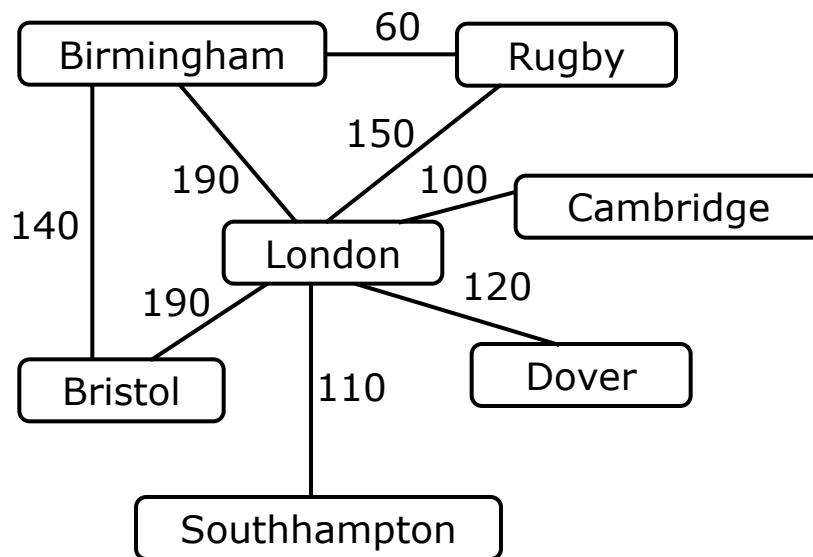
- A **graph** is a collection of **nodes** (or **vertices**, singular is **vertex**) and **edges** (or **arcs**)
 - Each node contains an **element**
 - Each edge connects two nodes together (or possibly the same node to itself) and may contain an **edge attribute**
- A **directed graph** is one in which the edges have a direction
- An **undirected graph** is one in which the edges do not have a direction
 - Note: Whether a graph is directed or undirected is a *logical* distinction—it describes how we think about the graph
 - Depending on the *implementation*, we may or may not be able to follow a directed edge in the “backwards” direction

Graph terminology II

- The **size** of a graph is the number of *nodes* in it
- The **empty graph** has size zero (no nodes)
- If two nodes are connected by an edge, they are **neighbors** (and the nodes are **adjacent** to each other)
- The **degree of a node** is the number of edges it has
- For directed graphs,
 - If a directed edge goes from node S to node D, we call S the **source** and D the **destination** of the edge
 - The edge is an **out-edge** of S and an **in-edge** of D
 - S is a **predecessor** of D, and D is a **successor** of S
 - The **in-degree** of a node is the number of in-edges it has
 - The **out-degree** of a node is the number of out-edges it has

Graph terminology III

- A **path** is a list of edges such that each node (but the last) is the predecessor of the next node in the list
- A **cycle** is a path whose first and last nodes are the same




- Example: (London, Bristol, Birmingham, London, Dover) is a path
- Example: (London, Bristol, Birmingham, London) is a cycle
- A **cyclic graph** contains at least one cycle
- An **acyclic graph** does not contain any cycles

Graph terminology IV

- An undirected graph is **connected** if there is a path from every node to every other node
- A *directed graph* is **strongly connected** if there is a path from every node to every other node
- A directed graph is **weakly connected** if the underlying undirected graph is connected
- Node **X** is **reachable** from node **Y** if there is a path from **Y** to **X**
- A subset of the nodes of the graph is a **connected component** (or just a **component**) if there is a path from every node in the subset to every other node in the subset


Adjacency-matrix representation I



	A	B	C	D	E	F	G
A				●	●		
B			●		●		
C		●					
D	●			●			●
E	●	●				●	●
F					●		
G				●	●		

- One simple way of representing a graph is the **adjacency matrix**
- A 2-D array has a mark at $[i][j]$ if there is an edge from node i to node j
- The adjacency matrix is symmetric about the main diagonal
- This representation is only suitable for *small* graphs! (Why?)

Adjacency-matrix representation II




	A	B	C	D	E	F	G
A					●		
B			●				
C							
D	●			●			
E		●					●
F					●		
G				●			

- An **adjacency matrix** can equally well be used for digraphs (directed graphs)
- A 2-D array has a mark at $[i][j]$ if there is an edge from node i to node j
- Again, this is only suitable for *small* graphs!

Edge-set representation I

- An **edge-set** representation uses a *set* of nodes and a *set* of edges
 - The sets might be represented by, say, linked lists
 - The set links are stored in the nodes and edges themselves
- The only other information in a node is its element (that is, its value)—it does not hold information about its edges
- The only other information in an edge is its source and destination (and attribute, if any)
 - If the graph is undirected, we keep links to both nodes, but don't distinguish between source and destination
- This representation makes it easy to find nodes from an edge, but you must search to find an edge from a node
- This is seldom a good representation

Edge-set representation II



nodeSet = {A, B, C, D, E, F, G}

edgeSet = { p: (A, E),
q: (B, C), r: (E, B),
s: (D, D), t: (D, A),
u: (E, G), v: (F, E),
w: (G, D) }


- Here we have a set of nodes, and each node contains only its element (not shown)

- Each edge contains references to its source and its destination (and its attribute, if any)

Adjacency-set representation I

- An adjacency-set representation uses a *set* of nodes
 - Each node contains a reference to the set of *its* edges
 - For a directed graph, a node might only know about (have references to) its out-edges
- Thus, there is not one single edge set, but rather a separate edge set for each node
 - Each edge would contain its attribute (if any) and its destination (and possibly its source)

Adjacency-set representation II




- Here we have a set of nodes, and each node refers to a set of edges
- Each edge contains references to its source and its destination (and its attribute, if any)

A	→	{ p }	p: (A, E)
B	→	{ q }	q: (B, C)
C	→	{ }	r: (E, B)
D	→	{ s, t }	s: (D, D)
E	→	{ r, u }	t: (D, A)
F	→	{ v }	u: (E, G)
G	→	{ w }	v: (F, E)
			w: (G, D)


Adjacency-set representation III

- If the edges have no associated attribute, there is no need for a separate **Edge** class
 - Instead, each node can refer to a set of its *neighbors*
 - In this representation, the edges would be implicit in the connections between nodes, not a separate data structure
- For an undirected graph, the node would have references to all the nodes adjacent to it
- For a directed graph, the node might have:
 - references to all the nodes adjacent to it, or
 - references to only those adjacent nodes connected by an out-edge from this node

Adjacency-set representation IV



- Here we have a set of nodes, and each node refers to a set of other (pointed to) nodes
- The edges are *implicit*



Searching a graph

- With certain modifications, any tree search technique can be applied to a graph
 - This includes depth-first, breadth-first, depth-first iterative deepening, and other types of searches
- The difference is that a graph may have cycles
 - We don't want to search around and around in a cycle
- To avoid getting caught in a cycle, we must keep track of which nodes we have already explored
- There are two basic techniques for this:
 - Keep a set of already explored nodes, or
 - Mark the node itself as having been explored
 - Marking nodes is not always possible (may not be allowed)

Example: Depth-first search

- **Here is how to do DFS on a tree:**

```
Put the root node on a stack;
while (stack is not empty) {
    remove a node from the stack;
    if (node is a goal node) return success;
    put all children of the node onto the stack;
}
return failure;
```

- **Here is how to do DFS on a graph:**

```
Put the starting node on a stack;
while (stack is not empty) {
    remove a node from the stack;
    if (node has already been visited) continue;
    if (node is a goal node) return success;
    put all adjacent nodes of the node onto the stack;
}
return failure;
```


Finding connected components

- A depth-first search can be used to find connected components of a graph
 - A connected component is a set of nodes; therefore,
 - A set of connected components is a set of sets of nodes
- To find the connected components of a graph:
 - while there is a node not assigned to a component {
 - put that node in a new component
 - do a DFS from the node, and put every node reached into the same component
 - }

Graph applications



- Graphs can be used for:
 - Finding a route to drive from one city to another
 - Finding connecting flights from one city to another
 - Determining least-cost highway connections
 - Designing optimal connections on a computer chip
 - Implementing automata
 - Implementing compilers
 - Doing garbage collection
 - Representing family histories
 - Doing similarity testing (e.g. for a dating service)
 - Pert charts
 - Playing games

Shortest-path

- Suppose we want to find the shortest path from node **X** to node **Y**
- It turns out that, in order to do this, we need to find the shortest path from **X** to *all* other nodes
 - Why?
 - If we don't know the shortest path from **X** to **Z**, we might overlook a shorter path from **X** to **Y** that contains **Z**
- **Dijkstra's Algorithm** finds the shortest path from a given node to *all* other reachable nodes

Dijkstra's algorithm I

- Dijkstra's algorithm builds up a *tree*: there is a path from each node back to the starting node
- For example, in the following graph, we want to find shortest paths from node **B**





- Edge values in the graph are weights
- Node values in the tree are *total* weights
- The arrows point in the *right direction* for what we need (why?)

Dijkstra's algorithm II

- For each vertex v , Dijkstra's algorithm keeps track of three pieces of information:
 - A boolean telling whether we *know* the shortest path to that node (initially true only for the starting node)
 - The length of the shortest path to that node known so far (0 for the starting node)
 - The predecessor of that node along the shortest known path (unknown for all nodes)
- Dijkstra's algorithm proceeds in phases—at each step:
 - From the vertices for which we don't know the shortest path, pick a vertex v with the smallest distance known so far
 - Set v 's “known” field to true
 - For each vertex w adjacent to v , test whether its distance so far is greater than v 's distance plus the distance from v to w ; if so, set w 's distance to the new distance and w 's predecessor to v

Dijkstra's algorithm III

- Three pieces of information for each node (e.g. **+3B**):
 - +** if the minimum distance is known *for sure*, blank otherwise
 - The best distance so far (**3** in the example)
 - The node's predecessor (**B** in the example, **-** for the starting node)




node	init'ly	<u>1</u>	<u>2</u>	<u>3</u>	<u>4</u>	<u>5</u>	<u>6</u>
A	inf	<u>3B</u>	3B	+3B	+3B	+3B	+3B
B	<u>0-</u>	+0-	+0-	+0-	+0-	+0-	+0-
C	inf	5B	<u>4A</u>	+4A	+4A	+4A	+4A
D	inf	inf	inf	<u>6C</u>	+6C	+6C	+6C
E	inf	inf	inf	8C	<u>+8C</u>	+8C	+8C
F	inf	inf	inf	inf	11D	<u>9E</u>	+9E

Summary

- A graph may be **directed** or **undirected**
- The **edges** (=arcs) may have weights or contain other data, or they may be just connectors
- Similarly, the **nodes** (=vertices) may or may not contain data
- There are various ways to represent graphs
 - The “best” representation depends on the problem to be solved
 - You need to consider what kind of access needs to be quick or easy
- Many tree algorithms can be modified for graphs
 - Basically, this means some way to recognize cycles

A graph puzzle



- Suppose you have a directed graph with the above shape
 - You don't know how many nodes are in the leader
 - You don't know how many nodes are in the loop
 - You don't know how many nodes there are total
 - You aren't allowed to mark nodes
- Devise an $O(n)$ algorithm (n being the total number of nodes) to decide when you *must already be* in the loop
 - This is *not* asking to find the *first* node in the loop
 - You can only use a *fixed* (constant) amount of extra memory

The End