


Queues

Stacks, Queues, and Deques

- A **stack** is a last in, first out (**LIFO**) data structure
 - Items are removed from a stack in the reverse order from the way they were inserted
- A **queue** is a first in, first out (**FIFO**) data structure
 - Items are removed from a queue in the same order as they were inserted
- A **deque** is a double-ended queue—items can be inserted and removed at either end

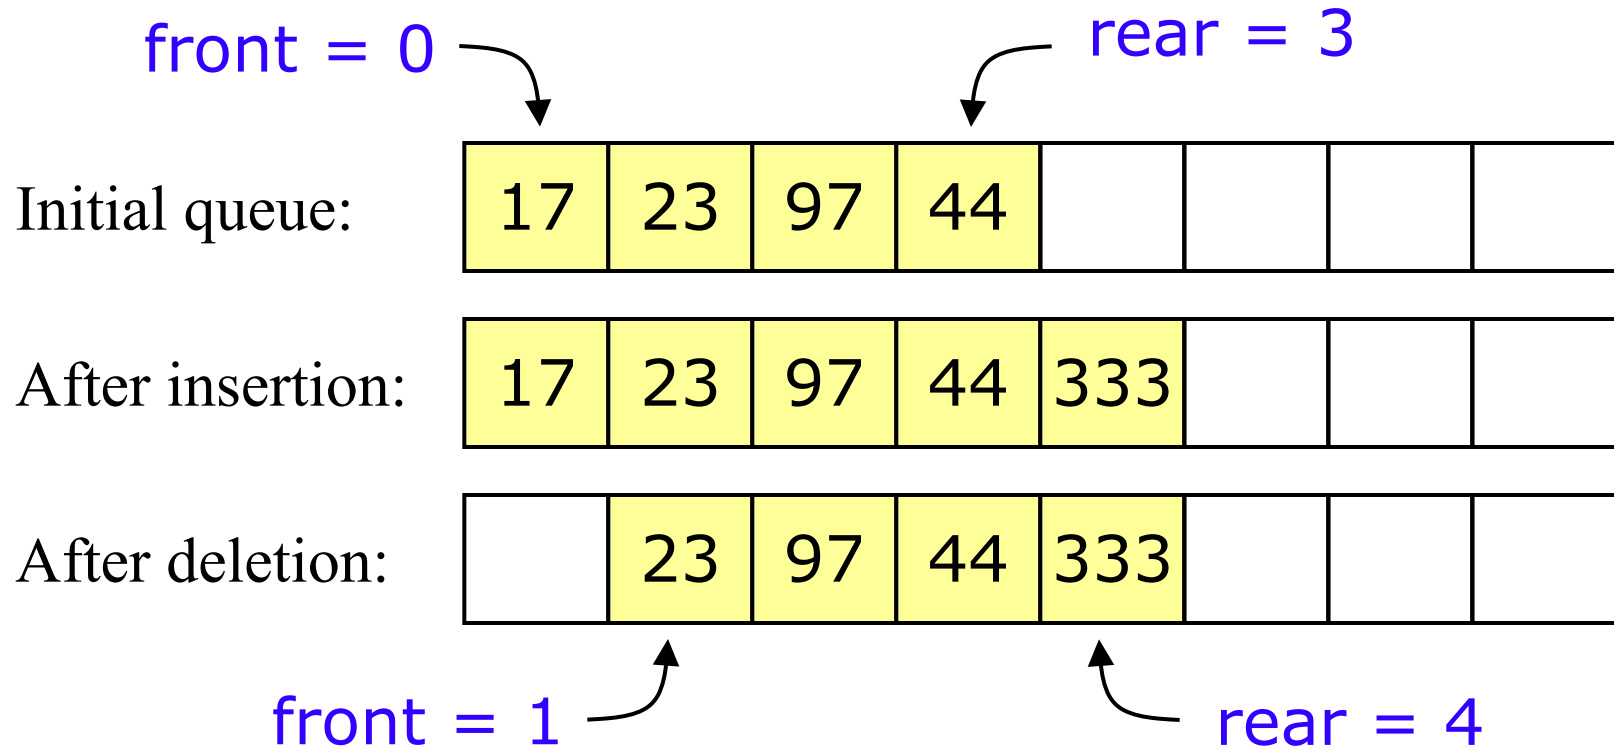
Array implementation of queues

- A **queue** is a first in, first out (**FIFO**) data structure
- This is accomplished by inserting at one end (the **rear**) and deleting from the other (the **front**)



- **To insert:** put new element in location 4, and set **rear** to 4
- **To delete:** take element from location 0, and set **front** to 1


Array implementation of queues



- Notice how the array contents “crawl” to the right as elements are inserted and deleted
- This will be a problem after a while!

Circular arrays


- We can treat the array holding the queue elements as circular (joined at the ends)




- Elements were added to this queue in the order 11, 22, 33, 44, 55, and will be removed in the same order
- Use: `front = (front + 1) % myQueue.length;`
and: `rear = (rear + 1) % myQueue.length;`

Full and empty queues

- If the queue were to become completely full, it would look like this:




- If we were then to remove all eight elements, making the queue completely empty, it would look like this:



This is a problem!

Full and empty queues: solutions

- **Solution:** Keep an additional variable




Linked-list implementation of queues

- In a queue, insertions occur at one end, deletions at the other end
- Operations at the front of a singly-linked list (SLL) are $O(1)$, but at the other end they are $O(n)$
 - Because you have to find the last element each time
- BUT: there is a simple way to use a singly-linked list to implement both insertions and deletions in $O(1)$ time
 - You always need a pointer to the first thing in the list
 - You can keep an additional pointer to the *last* thing in the list

SLL implementation of queues

- In an SLL you can easily find the successor of a node, but not its predecessor
 - Remember, pointers (references) are one-way
- If you know where the *last* node in a list is, it's hard to remove that node, but it's easy to add a node after it
- Hence,
 - Use the *first* element in an SLL as the *front* of the queue
 - Use the *last* element in an SLL as the *rear* of the queue
 - Keep pointers to *both* the front and the rear of the SLL

Enqueueing a node




To **enqueue** (add) a node:

- Find the current last node

- Change it to point to the new last node

- Change the **last** pointer in the list header

Dequeuing a node



- To **dequeue** (remove) a node:
 - Copy the pointer from the first node into the header

Queue implementation details

- With an array implementation:
 - you can have both overflow and underflow
 - you should set deleted elements to **null**
- With a linked-list implementation:
 - you can have underflow
 - overflow is a global out-of-memory condition
 - there is no reason to set deleted elements to **null**

A queue ADT

- Here are the usual operations on a queue:
 - `Queue()`: the constructor
 - `boolean isEmpty()`
 - `Object enqueue(Object item)`: add an element at the rear
 - `Object dequeue()`: remove an element from the front
 - `Object peek()`: look at the front element
 - `int search(Object o)`: Returns the 1-based position from the front of the queue

java.util Interface Queue<E>

- Java provides a queue *interface* and several implementations
- **boolean add(E e)**
 - Inserts the specified element into this queue if it is possible to do so immediately without violating capacity restrictions, returning true upon success and throwing an `IllegalStateException` if no space is currently available.
- **E element()**
 - Retrieves, but does not remove, the head of this queue.
- **boolean offer(E e)**
 - Inserts the specified element into this queue if it is possible to do so immediately without violating capacity restrictions.
- **E peek()**
 - Retrieves, but does not remove, the head of this queue, or returns null if this queue is empty.
- **E poll()**
 - Retrieves and removes the head of this queue, or returns null if this queue is empty.
- **E remove()**
 - Retrieves and removes the head of this queue.

Source: Java 6 API

A deque ADT

- Here are the operations expected of a deque:
 - `Deque()`: the constructor
 - `boolean isEmpty()`
 - `Object addAtFront(Object item)`
 - `Object addAtRear(Object item)`
 - `Object getFromFront()`
 - `Object getFromRear()`
 - `Object peekAtFront()`
 - `Object peekAtRear()`
 - `int search(Object o)`: Returns the 1-based position from the front of the deque

Using ArrayLists

- You could implement a deque with `java.util.ArrayList`:
 - `addAtFront(Object)` → `add(0, Object)`
 - `addAtRear(Object item)` → `add(Object)`
 - `getFromFront()` → `remove(0)`
 - `getFromRear()` → `remove(size() - 1)`
- If you did this, should you *extend* `ArrayList` or use it as a field in your `Deque` class?
- Would this be a good implementation?
- Why or why not?

java.util Interface Deque<E>

- Java 6 has a Deque interface
- There are 12 methods:
 - Add, remove, or examine an element...
 - ...at the head or the tail of the queue...
 - ...and either throw an exception, or return a special value (**null** or **false**) if the operation fails

	First Element (Head)		Last Element (Tail)	
	<i>Throws exception</i>	<i>Special value</i>	<i>Throws exception</i>	<i>Special value</i>
Insert	<u>addFirst(e)</u>	<u>offerFirst(e)</u>	<u>addLast(e)</u>	<u>offerLast(e)</u>
Remove	<u>removeFirst()</u>	<u>pollFirst()</u>	<u>removeLast()</u>	<u>pollLast()</u>
Examine	<u>getFirst()</u>	<u>peekFirst()</u>	<u>getLast()</u>	<u>peekLast()</u>

Priority queue

- A stack is first in, last out
- A queue is first in, first out
- A **priority queue** is *least-first-out*
 - The “smallest” element is the first one removed
 - (You could also define a *largest-first-out* priority queue)
 - The definition of “smallest” is up to the programmer (for example, you might define it by implementing **Comparator** or **Comparable**)
 - If there are several “smallest” elements, the implementer must decide which to remove first
 - Remove any “smallest” element (don’t care which)
 - Remove the first one added

A priority queue ADT

- Here is one possible ADT:
 - `PriorityQueue()`: a constructor
 - `void add(Comparable o)`: inserts `o` into the priority queue
 - `Comparable removeLeast()`: removes and returns the least element
 - `Comparable getLeast()`: returns (but does not remove) the least element
 - `boolean isEmpty()`: returns true iff empty
 - `int size()`: returns the number of elements
 - `void clear()`: discards all elements

Evaluating implementations

- When we choose a data structure, it is important to look at usage patterns
 - If we load an array once and do thousands of searches on it, we want to make searching fast—so we would probably sort the array
 - If we load a huge array and expect to do only a few searches, we probably *don't* want to spend time sorting the array
- For almost all uses of a queue (including a priority queue), we eventually remove everything that we add
- Hence, when we analyze a priority queue, neither “add” nor “remove” is more important—we need to look at the timing for “add + remove”

Array implementations

- A priority queue could be implemented as an *unsorted* array (with a count of elements)
 - Adding an element would take $O(1)$ time (why?)
 - Removing an element would take $O(n)$ time (why?)
 - Hence, adding *and* removing an element takes $O(n)$ time
 - This is an inefficient representation
- A priority queue could be implemented as a *sorted* array (again, with a count of elements)
 - Adding an element would take $O(n)$ time (why?)
 - Removing an element would take $O(1)$ time (why?)
 - Hence, adding *and* removing an element takes $O(n)$ time
 - Again, this is inefficient

Linked list implementations


- A priority queue could be implemented as an *unsorted* linked list
 - Adding an element would take $O(1)$ time (why?)
 - Removing an element would take $O(n)$ time (why?)
- A priority queue could be implemented as a *sorted* linked list
 - Adding an element would take $O(n)$ time (why?)
 - Removing an element would take $O(1)$ time (why?)
- As with array representations, adding *and* removing an element takes $O(n)$ time
 - Again, these are inefficient implementations

Binary tree implementations


- A priority queue could be represented as a (not necessarily balanced) binary search tree
 - Insertion times would range from $O(\log n)$ to $O(n)$ (why?)
 - Removal times would range from $O(\log n)$ to $O(n)$ (why?)
- A priority queue could be represented as a *balanced* binary search tree
 - Insertion and removal could destroy the balance
 - We need an algorithm to *rebalance* the binary tree
 - Good rebalancing algorithms require only $O(\log n)$ time, but are complicated

Heap implementation

- A priority queue can be implemented as a heap
- In order to do this, we have to define the *heap property*
 - In Heapsort, a node has the **heap property** if it is *at least as large* as its children
 - For a priority queue, we will define a node to have the **heap property** if it is *as least as small* as its children (since we are using smaller numbers to represent higher priorities)

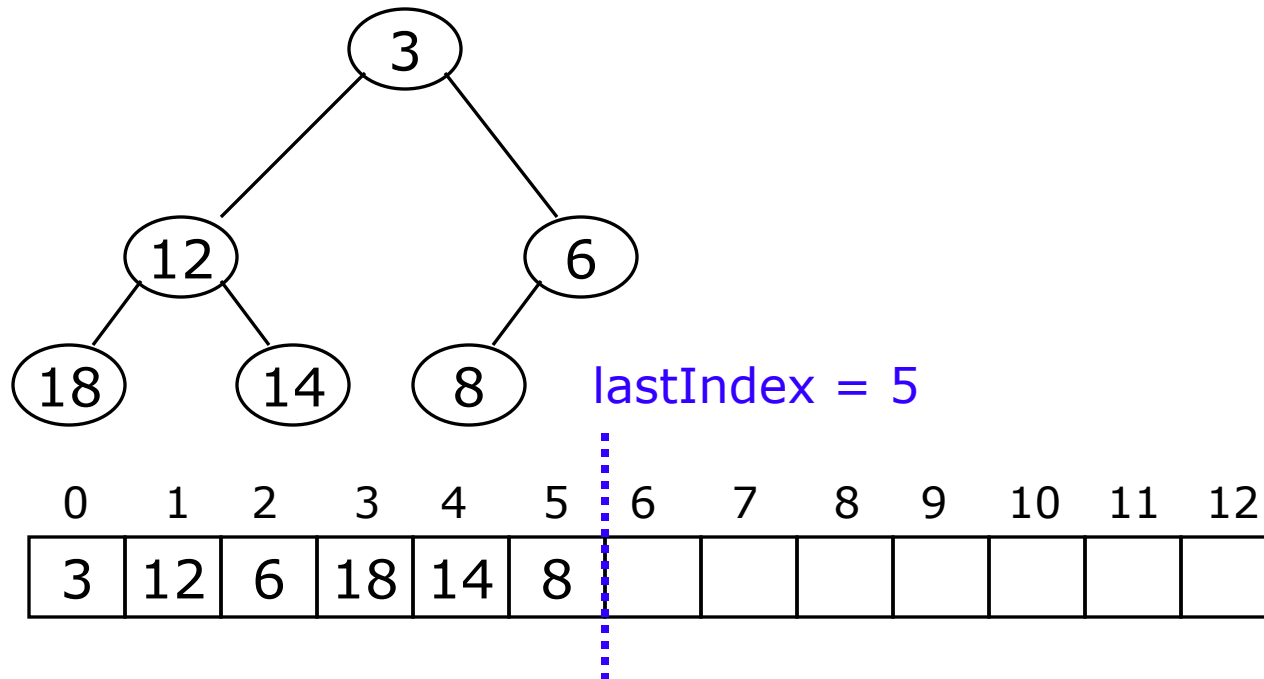


Heapsort: Blue node
has the heap property



Priority queue: Blue node
has the heap property

Array representation of a heap



- Left child of node i is $2*i + 1$, right child is $2*i + 2$
 - Unless the computation yields a value larger than lastIndex , in which case there is no such child
- Parent of node i is $(i - 1)/2$
 - Unless $i == 0$

Using the heap

- To add an element:
 - Increase `lastIndex` and put the new value there
 - Reheap the newly added node
 - This is called **up-heap bubbling**
 - Up-heap bubbling requires $O(\log n)$ time
- To remove an element:
 - Remove the element at location `0`
 - Move the element at location `lastIndex` to location `0`, and decrement `lastIndex`
 - Reheap the new root node (the one now at location `0`)
 - This is called **down-heap bubbling**
 - Down-heap bubbling requires $O(\log n)$ time
- Thus, it requires $O(\log n)$ time to add *and* remove an element

Comments

- A **priority queue** is a data structure that is designed to return elements in order of priority
- Efficiency is usually measured as the *sum* of the time it takes to add and to remove an element
 - Simple implementations take $O(n)$ time
 - Heap implementations take $O(\log n)$ time
 - Balanced binary tree implementations take $O(\log n)$ time
 - Binary tree implementations, without regard to balance, can take $O(n)$ (linear) time
- Thus, for any sort of heavy-duty use, heap or balanced binary tree implementations are better

Java 5 java.util.PriorityQueue

- Java 5 finally has a `PriorityQueue` class, based on heaps
 - Has redundant methods because it implements two similar interfaces
 - `PriorityQueue<E> queue = new PriorityQueue<E>();`
 - Uses the *natural ordering* of elements (that is, `Comparable`)
 - There is another constructor that takes a `Comparator` argument
 - `boolean add(E o)` – from the `Collection` interface
 - `boolean offer(E o)` – from the `Queue` interface
 - `E peek()` – from the `Queue` interface
 - `boolean remove(Object o)` – from the `Collection` interface
 - `E poll()` – from the `Queue` interface (returns `null` if queue is empty)
 - `void clear()`
 - `int size()`